

1. What is a support vector?

- Support vectors are the **data points that lie closest to the decision boundary (hyperplane)**.
 - They are critical because they directly define the position and orientation of the boundary.
 - Removing non-support vectors won't change the boundary, but removing a support vector will.
-

2. What does the C parameter do?

- **C is a regularization parameter** in SVM.
 - A **large C** → tries to classify all training points correctly (low bias, high variance, risk of overfitting).
 - A **small C** → allows more misclassifications but gives a simpler, smoother boundary (higher bias, lower variance).
-

3. What are kernels in SVM?

- Kernels are mathematical functions that let SVMs work in **higher-dimensional spaces** without explicitly computing the coordinates.
 - They measure similarity between data points.
 - Common kernels:
 - **Linear**: no transformation, just dot product.
 - **Polynomial**: compares points using polynomial features.
 - **RBF (Gaussian)**: compares points using distance similarity (localized).
-

4. Difference between Linear and RBF kernel?


- **Linear kernel**: assumes data is linearly separable (straight line/plane boundary).
 - **RBF kernel**: creates a flexible, non-linear boundary by mapping data into an infinite-dimensional feature space.
 - If features are already linearly separable → Linear is better.
 - If boundaries are complex/curved → RBF works better.
-

5. What are the advantages of SVM?

- Works well in **high-dimensional spaces**.
- Effective when number of features > number of samples.
- Robust to overfitting (especially with good regularization).

- Supports both **linear and non-linear classification** via kernels.
 - Often performs well even with limited data.
-

6. Can SVMs be used for regression?

-  Yes → called **Support Vector Regression (SVR)**.
 - Instead of finding a boundary, SVR fits a function that predicts continuous values within a tolerance margin (ϵ -insensitive loss).
-

7. What happens when data is not linearly separable?

- SVM introduces a **soft margin** (controlled by C) to allow misclassifications.
 - Or uses a **kernel trick** (like RBF) to transform data into higher dimensions where it may become separable.
-

8. How is overfitting handled in SVM?

- By tuning **C** (controls margin softness).
- By adjusting **kernel parameters** (e.g., gamma in RBF).
- By using **cross-validation** to pick hyperparameters.
- With **feature scaling/selection** to avoid noise dominance.